




Machine Learning Engineering & AI

Bootcamp Capstone

Step 2: Data Collection

Summary

Time Estimate: 3-6 Hours

To kick-start your Capstone Project, the first thing you'll need to do is collect your data. In some cases, it can be as simple as downloading a dataset in a ZIP file or a tarball. In other cases, this will require extracting data using a publicly available API or scraping a website. We urge you to work closely with your mentor to ensure that the data collection process is not too onerous for this Capstone Project. Also, **if your data collection requires you to write code, it's important that you start early.**

Below, you'll find a few great sources for large datasets that are appropriate for this course. These sources are also listed in some of the other project documentation.

- [Google dataset search](#)
- [AWS open datasets repository](#)
- [Uber Movement](#)
- [Yelp dataset](#)

In addition to the resources listed above, you can also explore datasets from [Quandl](#), [US Government Open Data](#), [UCI Machine Learning Repository](#), and [Kaggle competitions](#), or anywhere else you like.

Remember: Your dataset should have at least 15K-20K samples **at a minimum**. We'd like to see you build large-scale applications for this course so we encourage you to work with larger datasets, which means choosing a dataset that's at least 8GB in size or has at least 1 million samples.

Project Submission Steps

Please submit a link to a resource (e.g. GitHub repository), which should contain the following:

1. Three datasets that you have explored. Share with your mentors what make these datasets interesting to you. What is their prediction potential?
2. Code for how you collected the data, if applicable
3. Your actual chosen dataset: if your dataset is small enough to fit in a CSV, then feel free to include it in the repository. If it's a big dataset or has a lot of binary files like graphics or audio, consider linking to a cloud bucket (e.g. s3, Google Drive) in which you've stored your data. Alternatively, in the case of textual or tabular data, you may be able to check in the data in Git using [Git Large File Storage](#) extension. This is not recommended for truly large datasets.